

# **CSC 578**

## **Neural Networks and Deep Learning**

### **6-2. Convolutional Neural Networks (2)**

# What is an Image?



# Image

---



**An image is matrix that specifies the color of various pixels in terms of the amount of red, green and blue components.**


Every pixel is assigned an **RGB** value, each component being a value between 0 to 255. These components are mixed together to form a unique color value.

E.g. Blue is `rgb(0,0,255)` and Black is `rgb(0,0,0)`

## Displaying an Image



Any display of an image consists of an arrangement of red, green, and blue dots. A set of one dot of each color form a pixel





**What is a Kernel  Filter ?**

# Kernel



**A kernel is a square matrix that specifies spatial weights of various pixels in an image.**

Different image processing techniques use different kernels.

1	1	1
1	1	1
1	1	1

3\*3 Mean Kernel

1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1

5\*5 Gaussian Kernel

2	1	0
1	0	-1
0	-1	-2

3\*3 Sobel Kernel

# What is Convolution?

**BIG  
CONCEPT**





# Convolution



Convolution of a matrix involves laying a matrix over another and then calculating the weighted sum of all pixel values.

0	0	0	0	0	0
0	105	102	100	97	96
0	103	99	103	101	102
0	101	98	104	102	100
0	99	101	106	104	99
0	104	104	104	100	98

Image Matrix

Kernel Matrix		
0	-1	0
-1	5	-1
0	-1	0

320				

Output Matrix

$$\begin{aligned}
 &0 * 0 + 0 * -1 + 0 * 0 \\
 &+ 0 * -1 + 105 * 5 + 102 * -1 \\
 &+ 0 * 0 + 103 * -1 + 99 * 0 = 320
 \end{aligned}$$

Convolution with horizontal and  
vertical strides = 1




# Image Processing



# How to process an image?

---





RGB to Grayscale



Edge detection



-1	0	1
-2	0	2
-1	0	1





# Convolution Neural Network






# Convolution Layer




**32\*32\*3 image -> Preserve spatial structure**




# Convolution Layer



32x32x3 image

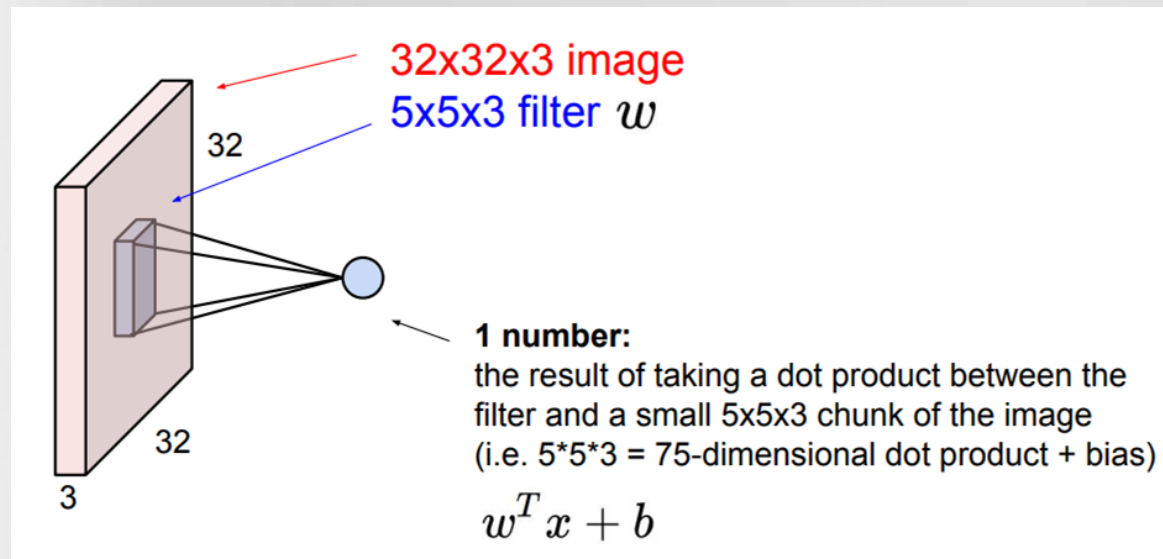


5x5x3 filter




**Convolve** the filter with the image  
i.e. “slide over the image spatially,  
computing dot products”

# Convolution Layer





# Convolution Layer






# Convolution Layer




Consider a second, **green** filter



# Convolution Layer



**E.g. if we had 6  $5 \times 5$  filters, we'll get 6 separate activation maps**




**We stack them up to get a “new image” of size  $28 \times 28 \times 6$**

# Convolution Layer



**ConvNet is a sequence of Convolution Layers, interspersed with activation functions**






# Convolution Layer



**ConvNet is a sequence of Convolution Layers, interspersed with activation functions**






# Fully Connected Layer



**32\*32\*3 image -> stretch to 3072\*1**





one filter =>  
one activation map

example 5x5 filters  
(32 total)

Activations:





Figure copyright Andrej Karpathy.

We call the layer convolutional  
because it is related to convolution  
of two signals:

$$f[x,y] * g[x,y] = \sum_{n_1=-\infty}^{\infty} \sum_{n_2=-\infty}^{\infty} f[n_1,n_2] \cdot g[x-n_1,y-n_2]$$

↑  
elementwise multiplication and sum of  
a filter and the signal (image)

## A closer look at spatial dimension (1)







7


7

7x7 input (spatially)  
assume 3x3 filter





7


7

7x7 input (spatially)  
assume 3x3 filter



7


7

7x7 input (spatially)  
assume 3x3 filter



7


7

7x7 input (spatially)  
assume 3x3 filter






7


7

7x7 input (spatially)  
assume 3x3 filter


**=> 5x5 output**

## A closer look at spatial dimension



7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 2**

## A closer look at spatial dimension (2)



7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 2**






7



7

7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 2**  
**=> 3x3 output!**

## A closer look at spatial dimension (3)



7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 3?**



7x7 input (spatially)  
assume 3x3 filter  
applied **with stride 3?**

**doesn't fit!**  
cannot apply 3x3 filter on  
7x7 input with stride 3.



## Constraints on Output Size



For the input volume size (W), the receptive field size of the Conv Layer neurons (F), the stride with which they are applied (S), and the amount of zero padding used (P) on the border, the size of the output is

$$(W-F+2P)/S+1$$

### Examples:

- For a 7x7 input and a 3x3 filter with stride 1 and pad 0, we would get a 5x5 output.
- With the same input and filter but with stride 2, we would get a 3x3 output.



## Use of zero-padding

- To avoid shrinking output
- To throw away information from edges

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

**3x3** filter, applied with **stride 1**

**pad with 1 pixel** border => what is the output?

**7x7 output!**


**Setting zero padding to be  $P=(F-1)/2$  when the stride  $S=1$  ensures that the input volume and output volume will have the same size spatially.**



## Exercise (1)

Input volume: **32x32x3**  
10 5x5 filters with stride 1, pad 2

Output volume size: ?








Input volume: **32x32x3**

**10** **5x5** filters with stride **1**, pad **2**

Output volume size:

$(32 + 2 * 2 - 5) / 1 + 1 = 32$  spatially, so

**32x32x10**




## Exercise (2)



Input volume: **32x32x3**


10 5x5 filters with stride 1, pad 2

Number of parameters in this layer?





Input volume: **32x32x3**  
**10** **5x5** filters with stride 1, pad 2



Number of parameters in this layer?  
each filter has  $5*5*3 + 1 = 76$  params (+1 for bias)  
 $\Rightarrow 76*10 = 760$


**Reference:** <http://cs231n.github.io/convolutional-networks/>  
for explanation and animation on the DEPTH of input  
volume.



## CAUTION:



**E.g.  $32 \times 32$  input convolved repeatedly with  $5 \times 5$  filters shrinks volumes spatially.**




40

**BUT shrinking too fast is not good, because feature maps lose information as they get deeper.**

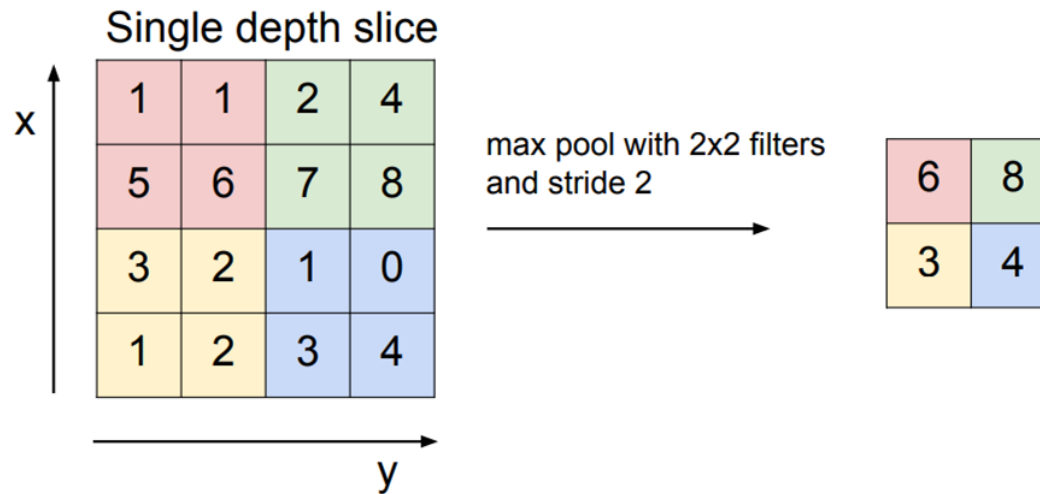
## Pooling Layer



**Makes the representations smaller and more manageable**




# Max Pooling





## Intervening Activation Layer



# CNN Architectures



# AlexNet



[Krizhevsky et al. 2012]

## Architecture:

CONV1

MAX POOL1

NORM1

CONV2

MAX POOL2

NORM2

CONV3

CONV4

CONV5

Max POOL3

FC6

FC7

FC8





Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.



# AlexNet




Input: 227x227x3 images

**First layer (CONV1):** 96 11x11 filters applied at stride 4

=>

Q: what is the output volume size? Hint:  $(227-11)/4+1 = 55$

# AlexNet



Input: 227x227x3 images


**First layer (CONV1):** 96 11x11 filters applied at stride 4

=>

Output volume **[55x55x96]**

Q: What is the total number of parameters in this layer?

# AlexNet



Input: 227x227x3 images

**First layer (CONV1):** 96 11x11 filters applied at stride 4


=>

Output volume **[55x55x96]**

Parameters:  $(11*11*3)*96 = 35K$



# AlexNet




Input: 227x227x3 images

After CONV1: 55x55x96

**Second layer (POOL1):** 3x3 filters applied at stride 2

Q: what is the output volume size? Hint:  $(55-3)/2+1 = 27$

# AlexNet




Input: 227x227x3 images  
After CONV1: 55x55x96

**Second layer (POOL1):** 3x3 filters applied at stride 2  
Output volume: 27x27x96

Q: what is the number of parameters in this layer?

# AlexNet



Input: 227x227x3 images  
After CONV1: 55x55x96

**Second layer (POOL1):** 3x3 filters applied at stride 2  
Output volume: 27x27x96  
Parameters: 0!



# VGGNet



# GoogleNet

